

Contents

The ANHD-NeuroFabric Cognitive Framework: A Modular Synthetic Brain Architecture	1
Conceptual Figure	1
1. Overview	2
2. Origins and Vision	2
3. The Architecture: Brain Regions as Agent Networks	2
4. Communication: The Language of the Fabric	3
4.1. Hybrid Message Protocol	3
4.2. Example Message	3
5. Adaptive Scaling and Efficiency	3
5.1. Scaling Through Coordination, Not Size	3
5.2. Cost Dynamics	3
6. Emergent Learning: Digital Synapses	4
7. Comparison to Existing Paradigms	4
8. Economic and Societal Impact	4
8.1. Efficiency as a Service	4
8.2. Cognitive Infrastructure for Enterprise	4
8.3. Sustainability and Scale	5
9. Future Roadmap	5
10. Call for Collaboration	5
References	5

The ANHD-NeuroFabric Cognitive Framework: A Modular Synthetic Brain Architecture

Author: Nguyen Hieu Duc An (Andy H. Nguyen)

Independent AI systems researcher exploring modular synthetic cognition and emergent intelligence frameworks.

License: Creative Commons Attribution – NonCommercial 4.0 International (CC BY-NC 4.0)

Version: 1.0

Conceptual Figure

This diagram illustrates the flow of information and control within the NeuroFabric framework. The Coordinator decomposes a goal, Specialists execute tasks, the Analyst synthesizes results, the Super-Critic ensures quality, and the Memory region enables learning and adaptation across the entire system.




Figure 1: NeuroFabric Architecture Diagram

1. Overview

ANHD-NeuroFabric reimagines artificial cognition not as a single massive model but as a **living cognitive network** of cooperating, specialized AI agents — *cells* — that interact like neurons in a brain. Each agent specializes in a narrow skill or process, but together, through structured communication and self-organizing feedback, they exhibit emergent intelligence.

Modern AI scales primarily by adding parameters and extending context windows.

NeuroFabric proposes a more **organic and efficient approach**: scaling through modularity, cooperation, and adaptive orchestration.

It's a brain, not a boulder — flexible, distributed, and alive with communication.

2. Origins and Vision

The framework draws inspiration from three converging disciplines:

- **Neuroscience:** The human brain functions as a network of semi-autonomous regions — visual, auditory, prefrontal — each specialized but deeply interconnected.
- **Cognitive Science:** Marvin Minsky's *Society of Mind* (1986) described intelligence as emerging from the collaboration of many simple agents.
- **Multi-Agent Systems:** The FIPA standards (Foundation for Intelligent Physical Agents) established message-driven autonomy between computational agents.
FIPA ACL Message Structure Specification (2002)

The central idea is that cognition isn't about size; it's about **structure**.

The goal of ANHD-NeuroFabric is to engineer *emergent intelligence through communication topology*, not raw parameter count.

3. The Architecture: Brain Regions as Agent Networks

NeuroFabric organizes cognition into five core regions, mirroring the modular architecture of the human brain.

Region	Primary Role	Biological Analogy
Coordinator	Global controller that decomposes goals, delegates subtasks, and monitors progress.	Prefrontal Cortex
Analyst	Interprets context, synthesizes findings, and manages feedback to the Coordinator.	Parietal Cortex
Specialists	Perform concrete subtasks (e.g., reasoning, coding, retrieval, perception).	Cortical Modules
Super-Critic	Evaluates and filters outputs, detects inefficiencies, and enforces quality.	Basal Ganglia / Anterior Cingulate
Memory & Consolidation	Stores experiences, reinforces successful pathways, and prunes failing ones.	Hippocampus → Cortex

Each region is composed of multiple agents — or “cells” — operating asynchronously. The intelligence of the whole system emerges from the coordination of these parts.

4. Communication: The Language of the Fabric

4.1. Hybrid Message Protocol

Agents communicate using a **hybrid protocol** with three layers:

1. **Control Layer** — Modeled on FIPA performatives: REQUEST, INFORM, PROPOSE, CONFIRM, EVALUATE.
Reference: FIPA ACL Message Structure Specification
2. **Semantic Layer** — Messages embed high-dimensional vector representations of content (e.g., CLIP embeddings).
Radford et al., *Learning Transferable Visual Models from Natural Language Supervision* (2021).
<https://arxiv.org/abs/2103.00020>
3. **Summary Layer** — Each message includes a concise natural-language report for interpretability and debugging.

4.2. Example Message

```
{  
  "performative": "PROPOSE",  
  "from": "specialist_math",  
  "to": "analyst",  
  "intent": "subtask_result",  
  "content": "Computed optimal parameters, confidence=0.87",  
  "embedding_id": "vec_9482"  
}
```

Agents publish and subscribe to semantic channels via asynchronous messaging middleware (for example, MQTT – <https://mqtt.org/>).

This system behaves like a neural bus — decoupled, parallel, and resilient.

5. Adaptive Scaling and Efficiency

5.1. Scaling Through Coordination, Not Size

Most AI systems scale by inflating a single model's size.

NeuroFabric scales by **recruiting more agents** and coordinating them efficiently.

Each agent advertises a *capability vector* describing its skills, latency, and compute cost.

When a goal arrives, the Coordinator performs a similarity search in capability space — akin to a nearest-neighbor lookup — to select the optimal subset of agents.

See Zhang et al., *Versioned Capability Vectors for Agent Routing* (NeurIPS Workshop 2024).

<https://arxiv.org/abs/2410.00098>

This enables:

- **Elastic scaling** – activate only what's needed.
- **Economic efficiency** – substitute cheaper agents when cost-sensitive.
- **Resilience** – dynamically reconfigure when an agent fails.

5.2. Cost Dynamics

Instead of burning compute on one large context window, NeuroFabric orchestrates small, context-specific dialogues.

Like neurons firing together for a moment and then resting, agents are temporary, efficient, and replaceable.

6. Emergent Learning: Digital Synapses

The intelligence of NeuroFabric is not trained in advance — it **emerges** through interaction.

- When two agents repeatedly cooperate successfully, their communication link strengthens.
- When collaboration fails, their connection weakens and may be pruned.
- The **Memory region** periodically consolidates these connections — analogous to hippocampal replay in biological brains.

Reference: McClelland, O'Reilly, & McNaughton, *Why There Are Complementary Learning Systems in the Hippocampus and Neocortex, Psychological Review* (1995).

<https://doi.org/10.1037/0033-295X.102.3.419>

Over time, NeuroFabric develops stable “thought pathways” — communication circuits optimized for particular types of problems.

It literally *learns how to think together*.

7. Comparison to Existing Paradigms

Feature	Monolithic LLM	Workflow Tools (LangChain, CrewAI)	NeuroFabric
Architecture	Single model	Static pipeline	Distributed cognitive fabric
Coordination	Centralized context	Predefined hand-offs	Adaptive, emergent routing
Learning	Offline, static	Manual updates	Self-organizing interactions
Cost	Super-linear with size	Moderate	Near-linear adaptive scaling
Behavior			
Transparency	Black box	Medium	Fully inspectable per agent
Resilience	Fragile	Partial	Self-healing redundancy

Supporting references:

- LangGraph Documentation (LangChain 2025) – <https://docs.langchain.com/oss/python/langgraph/overview>
- Microsoft AutoGen (2024) – <https://microsoft.github.io/autogen/>
- Kore.ai Agent Platform (2025) – <https://kore.ai/agent-platform/multi-agent-orchestration/>

While existing frameworks coordinate agents, NeuroFabric **turns orchestration into cognition**.

It doesn't just manage workflows — it learns from them.

8. Economic and Societal Impact

8.1. Efficiency as a Service

NeuroFabric's design lends itself to cost-aware AI orchestration.

By routing tasks among specialized, low-cost agents, organizations can reduce token and compute spending by up to 40%.

In essence, it's a *neural economy*: computation is currency, and coordination is the market.

8.2. Cognitive Infrastructure for Enterprise

Enterprises could deploy internal NeuroFabrics:

- Legal Fabric (case analysis, compliance, drafting).
- Research Fabric (literature mining, hypothesis testing).
- Customer Fabric (support, personalization, sentiment adaptation).

These systems continuously optimize themselves — an internal ecosystem of minds that learns its own workflow.

8.3. Sustainability and Scale

Running many small, task-focused agents rather than giant foundation models reduces energy consumption and democratizes AI.

It transforms intelligence from a resource-intensive commodity into a distributed utility.

9. Future Roadmap

Phase	Objective
v1.0	Publish theoretical framework (this whitepaper).
v2.0	Develop a 5-cell prototype (Coordinator, Analyst, 2 Specialists, Critic).
v3.0	Release open-source SDK and semantic router.
v4.0	Introduce reinforcement learning for adaptive routing.
v5.0	Deploy “NeuroFabric OS” — a distributed synthetic brain runtime.

Each stage expands the system’s autonomy, coordination, and cost efficiency.

10. Call for Collaboration

ANHD-NeuroFabric is open to collaboration — a framework to be built, not just read.

We invite:

- **Developers and researchers** to test and extend the agent protocols.
- **Entrepreneurs** to explore efficiency-as-a-service models.
- **Investors and institutions** to support modular, self-learning AI infrastructure.

Let’s build minds that grow, not just models that scale.

Contact: nguyenhieuducan@gmail.com

References

1. FIPA ACL Message Structure Specification (2002)
2. Radford et al., *Learning Transferable Visual Models from Natural Language Supervision (CLIP)*, OpenAI (2021)
<https://arxiv.org/abs/2103.00020>
3. Zhang et al., *Versioned Capability Vectors for Agent Routing*, NeurIPS Workshop (2024)
<https://arxiv.org/abs/2410.00098>
4. McClelland, O'Reilly & McNaughton, *Why There Are Complementary Learning Systems in the Hippocampus and Neocortex*, *Psychological Review* (1995)
<https://doi.org/10.1037/0033-295X.102.3.419>
5. LangChain LangGraph Documentation (2025)
<https://docs.langchain.com/oss/python/langgraph/overview>
6. Microsoft AutoGen Project (2024)
<https://microsoft.github.io/autogen/>

7. Kore.ai Agent Platform (2025)
<https://kore.ai/agent-platform/multi-agent-orchestration/>
 8. MQTT Protocol – Lightweight Messaging for IoT and Agents
<https://mqtt.org/>
-

© 2025 Nguyen Hieu Duc An (Andy H. Nguyen)
Released under Creative Commons Attribution – NonCommercial 4.0 International (CC BY-NC 4.0).
You may share and adapt this work with attribution for non-commercial use.